



Parallel Programming with Threads and Tasks

TDDD56 Lecture 2

Christoph Kessler

PELAB / IDA
Linköping university
Sweden

2014



Outline

Lecture 1: Multicore Architecture Concepts

- Architectural trends and consequences for programming

Lecture 2: Parallel programming with threads and tasks

- Revisiting processes, threads, synchronization
 - Pthreads
 - OpenMP (very shortly)
- Tasks
 - Cilk
 - Futures

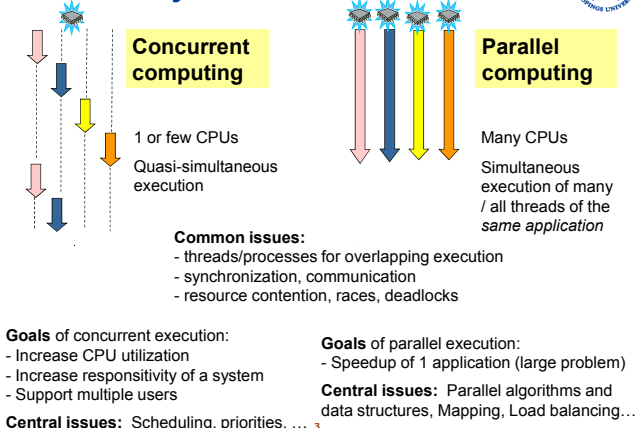
Lecture 3: Shared memory architecture concepts

Lecture 4: Non-blocking synchronization

Lecture 5: Design and analysis of parallel algorithms



Concurrency vs. Parallelism




Processes

(Refresher from TDD68)



Example: Process Creation in UNIX

- **fork** system call
 - creates new child process
- **exec** system call
 - used after a **fork** to replace the process' memory space with a new program
- **wait** system call
 - by parent, suspends parent execution until child process has terminated



```
int main()
{
    Pid_t ret;
    /* fork another process: */
    ret = fork();
    if (ret < 0) { /* error occurred */
        fprintf ( stderr, "Fork Failed" );
        exit(-1);
    }
    else if (ret == 0) { /* child process */
        execlp ( "/bin/ls", "ls", NULL );
    }
    else { /*parent process: ret=childPID
          /* will wait for child to complete: */
        wait (NULL);
        printf ("Child Complete");
        exit(0);
    }
}
```

5

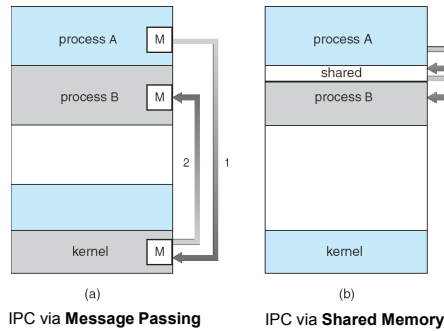


Parallel programming with processes

- Processes can create new processes that execute concurrently with the parent process
- OS scheduler – also for single-core CPUs
- Different processes share nothing by default
 - Inter-process communication via OS only, via shared memory (write/read) or message passing (send/recv)
- **Threads** are a more light-weight alternative for programming shared-memory applications
 - Sharing memory (except local stack) by default
 - Lower overhead for creation and scheduling/dispatch
 - ▶ E.g. Solaris: creation 30x, switching 5x faster

6

IPC Models – Realization by OS

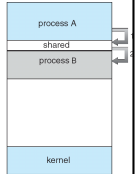


Syscalls: send, rcv

Syscalls: shmget, shmat, then load / store

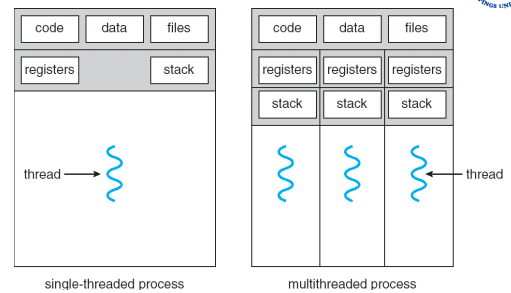
Example: POSIX Shared Memory API

- `#include <sys/shm.h>`
`#include <sys/stat.h>`
- Let OS create a shared memory segment (system call):
 - `int segment_id = shmget (IPC_PRIVATE, size, S_IRUSR | S_IWUSR);`
- Attach the segment to the executing process (system call):
 - `void *shmemptr = shmat (segment_id, NULL, 0);`
- Now access it:
 - `strcpy ((char *)shmemptr, "Hello world");` // Example: copy a string into it
 - ...
- Detach it from executing process when no longer accessed:
 - `shmdt (shmemptr);`
- Let OS delete it when no longer used:
 - `shmctl (segment_id, IPC_RMID, 0);`



Threads

Single- and Multithreaded Processes



A **thread** is a basic unit of CPU utilization:

- Thread ID, program counter, register set, stack.

A process may have one or several threads.

TDD868, C. Kessler, IDA, Linköpings universitet.

4.10

Silberschatz, Galvin and Gagne ©2005

Benefits of Multithreading

- Responsiveness
 - Interactive application can continue even when part of it is blocked
- Resource Sharing
 - Threads of a process share its memory by default.
- Economy
 - Light-weight
 - Creation, management, context switching for threads is much faster than for processes
- Utilization of Multiprocessor Architectures
 - Convenient (but low-level) shared memory programming

TDD868, C. Kessler, IDA, Linköpings universitet.

4.11

POSIX Threads (Pthreads)

- A POSIX standard (IEEE 1003.1c) API for thread programming in C
 - start and terminate threads
 - coordinate threads
 - regulate access to shared data structures
- API specifies behavior, not implementation, of the thread library
- C interface, e.g.
 - `int pthread_create (pthread_t *thread, const pthread_attr_t *attr, void *(*start_routine)(void*), void *arg);`
- Note: as a library, rely on underlying OS and hardware!
- Common in UNIX operating systems (Solaris, Linux, Mac OS X)

C. Kessler, IDA, Linköpings universitet.

12

Starting a Thread (1)

- Thread is started with function

```
int pthread_create ( pthread_t *thread, const pthread_attr_t *attr,
void *(*func)(void*), void *arg);
```

- Called func must have parameter and ret values void*
 - Exception: first thread is started with main()
- Thread terminates when called function terminates or by pthread_exit (void *retval)
- Threads started one by one
- Threads represented by data structure of type pthread_t

C. Kessler, IDA, Linköping universitet.

13

Starting a Thread (2)

- Example:

```
#include <pthread.h>

int main ( int argc, char *argv[] )
{
    int *ptr;
    pthread_t thr;

    pthread_create( &thr,
                    NULL,
                    foo,
                    (void*)ptr );

    ...
    pthread_join( &thr, NULL );
    return 0;
}
```

```
void *foo ( void *vp )
{
    int i = (int) vp;;
    ...

    // alternative
    // – pass a parameter block:
    void *foo ( void *vp )
    {
        Userdefinedstructtype *ptr;
        ptr=(Userdefinedstructtype*)vp;
        ...
    }
}
```

C. Kessler, IDA, Linköping universitet.

14

Access to Shared Data (0)

- Globally defined variables are globally shared and visible to all threads.
- Locally defined variables are visible to the thread executing the function.
- But all data in shared memory publish an address of data: all threads could access...
- Take care: typically no protection between thread data – thread1 (foo1) could even write to thread2's (foo2) stack frame

- Example 0: Parallel incrementing

```
int a[N]; // shared, assume P | N
pthread_t thr[P];

int main( void )
{
    int t;
    for (t=0; t<P; t++)
        pthread_create(&(thr[t]), NULL,
                        incr, a + t*N/P );
    for (t=0; t<P; t++)
        pthread_join( thr[t], NULL );
    ...
    void *incr ( void *myptr_a )
    {
        int i;
        for (i=0; i<N/P; i++)
            ((int*)myptr_a[i])++;
    }
}
```

C. Kessler, IDA, Linköping universitet.

Access to Shared Data (1)

- Globally defined variables are globally shared and visible to all threads.
- Locally defined variables are visible to the thread executing the function.
- But all data in shared memory publish an address of data: all threads could access...
- Take care: typically no protection between thread data – thread1 (foo1) could even write to thread2's (foo2) stack frame

- Example 1

```
int *globalptr = NULL; // shared ptr

void *foo1 ( void *ptr1 )
{
    int i = 15;
    globalptr = &i; // ??? dangerous!
    // if foo1 terminates, foo2 writes
    // somewhere, unless globalptr
    // value is reset to NULL manually
    ...
}

void *foo2 ( void *ptr2 )
{
    if (globalptr) *globalptr = 17;
    ...
}
```

C. Kessler, IDA, Linköping universitet.

Access to Shared Data (2)

- Globally defined variables are globally shared and visible to all threads
- Locally defined variables are visible to the thread executing the function
- But all data in shared memory publish an address of data: all threads could access...
- Take care: typically no protection between thread data – thread1 could even write to thread2's stack frame

- Example 2

```
int *globalptr = NULL; // shared ptr

void *foo1 ( void *ptr1 )
{
    int i = 15;
    globalptr =(int*)malloc(sizeof(int));
    // safe, but possibly memory leak;
    // OK if garbage collection ok
    ...
}

void *foo2 ( void *ptr2 )
{
    if (globalptr) *globalptr = 17;
    ...
}
```

C. Kessler, IDA, Linköping universitet.

Coordinating Shared Access (3)

What if several threads need to write a shared variable?

- If they simply write: ok if write order does not play a role
- If they read and write: encapsulate (critical section, monitor) and protect e.g. by mutual exclusion using mutex locks
- Example: Access to a taskpool
 - threads maintain list of tasks to be performed
 - if thread is idle, gets a task and performs it

```
// each thread:
while (!workdone)
{
    task = gettask( Pooldescr );
    performtask ( task );
}

// may be called concurrently:
Tasktype gettask ( Pool p )
{
    // begin critical section
    task = p.queue [ p.index ];
    p.index++;
    // end critical section

    return task;
}
```

C. Kessler, IDA, Linköping universitet.


18

Race Conditions lead to Nondeterminism



- Example: `p.index++`
- could be implemented in machine code as

```
39: register1 = p.index    // load
40: register1 = register1 + 1 // add
41: p.index = register1    // store
```



- Consider this execution interleaving, with "index = 5" initially:

39: thread1 executes register1 = p.index	{ T1.register1 = 5 }
39: thread2 executes register1 = p.index	{ T2.register1 = 5 }
40: thread1 executes register1 = register1 + 1	{ T1.register1 = 6 }
40: thread2 executes register1 = register1 + 1	{ T2.register1 = 6 }
41: thread1 executes p.index = register1	{ p.index = 6 }
41: thread2 executes p.index = register1	{ p.index = 6 }

- Compare to a different interleaving, e.g., 39,40,41, 39,40,41...
- Result depends on relative speed of the accessing threads (*race condition*)

C. Kessler, IDA, Linköping universitet.

19

Critical Section



- **Critical Section:** A set of instructions, operating on shared data or resources, that should be executed by a single thread at a time without interruption
- **Atomicity** of execution
- **Mutual exclusion:** At most one process should be allowed to operate inside at any time
- **Consistency:** inconsistent intermediate states of shared data not visible to other processes outside
- May consist of different program parts for different threads
 - that access the same shared data



- General structure, with structured control flow:

```
...
Entry of critical section C
... critical section C: operation on shared data
Exit of critical section C
```

C. Kessler, IDA, Linköping universitet.

20

Coordinating Shared Access (4)



```
pthread_mutex_t mutex; // global variable definition - shared
```

```
...
// in main:
```

```
pthread_mutex_init( &mutex, NULL );
```

```
...
// in gettask:
```

```
pthread_mutex_lock( &mutex );
```

```
task = p.queue [p.index];
```

```
p.index++;
```

```
pthread_mutex_unlock( &mutex );
```

```
...
```

Often implemented using test_and_set or other atomic instruction where available

C. Kessler, IDA, Linköping universitet.

21

Hardware Support for Synchronization



- Most systems provide hardware support for protecting critical sections
- Uniprocessors – could **disable interrupts**
 - Currently running code would execute without preemption
 - Generally too inefficient on multiprocessor systems
 - ▶ Operating systems using this are not broadly scalable
- Modern machines provide special **atomic instructions**
 - **TestAndSet:** test memory word and set value atomically
 - ▶ Atomic = non-interruptable
 - ▶ If multiple TestAndSet instructions are executed *simultaneously* (each on a different CPU in a multiprocessor), then they take effect sequentially in some arbitrary order.
 - **AtomicSwap:** swap contents of two memory words atomically
 - **CompareAndSwap**
 - **Load-linked / Store-conditional**

C. Kessler, IDA, Linköping universitet.

22

TestAndSet Instruction



- Definition in pseudocode:

```
boolean TestAndSet (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv; // return the OLD value
}
```



C. Kessler, IDA, Linköping universitet.

23

Mutual Exclusion using TestAndSet



- Shared boolean variable **lock**, initialized to FALSE (= unlocked)
- **do {**
 - while (TestAndSet (&lock))**
 - ; // do nothing but spinning on the lock (busy waiting)
 - // ... critical section
 - lock = FALSE;
 - // ... remainder section
- } while (TRUE);**

C. Kessler, IDA, Linköping universitet.

24

→EXAMPLE

Compare-And-Swap

see Lecture 4

- CompareAndSwap (adr_memcell, value, register)
 - Atomically compares a value in a memory cell to a supplied value and, if these are equal, swaps the contents of the memory cell with the value stored in a register.

■ Example: Mutual Exclusion using Compare-And-Swap:

```

register = 1;
CompareAndSwap ( &lock, 0, register );
while (register != 0)
    CompareAndSwap ( &lock, 0, register );
// ... critical section
lock = 0;
// ...
    
```

- More details in the lecture on non-blocking synchronization

C. Kessler, IDA, Linköping universitet.

25

Load-linked / Store-conditional

see Lecture 4

2 new instructions for memory access:

- **LoadLinked** address, register
 - records the version number of the value read (cf. a svn update)
- **StoreConditional** register, address
 - will only succeed if no other operations were executed on the accessed memory location since my last LoadLinked instruction to address, (cf. a svn commit)
 - and set the register operand of Store-conditional to 0, otherwise.

C. Kessler, IDA, Linköping universitet.

26

Mutual Exclusion using LL/SC

see Lecture 4

- Shared int variable lock, initialized to 0 (= unlocked)

```

do {
    register = 0;
    while ( register == 0 ) {
        dummy = LoadLinked ( &lock );
        if (dummy == 0) { // read a 0 – found unlocked
            register = 1;
            register = StoreConditional ( register, &lock );
            // if register is 0, StoreConditional failed, retry...
        }
    }
    // ... critical section
    lock = 0; // ordinary store
    // ... remainder section
} while ( TRUE);
    
```

C. Kessler, IDA, Linköping universitet.

27

Pitfalls with Semaphores

- Correct use of mutex operations:

- Protect all possible entries/exits of control flow into/from critical section:

```

pthread_mutex_lock (&mutex)
....
pthread_mutex_unlock (&mutex)
    
```



- Possible sources of synchronization errors:

- Omitting lock(&mutex) or unlock(&mutex) (or both) ??
- lock(&mutex) lock(&mutex) ??
- lock(&mutex1) unlock(&mutex2) ??
- if-statement in critical section, unlock in then-branch only

C. Kessler, IDA, Linköping universitet.

28

Problems: Deadlock and Starvation

- **Deadlock** – two or more threads are waiting indefinitely for an event that can be caused only by one of the waiting threads

- Typical example: *Nested critical sections*

- ▶ Guarded by locks S and Q, initialized to unlocked

	P_0	P_1
time ↓	wait (S);	wait (Q);
	wait (Q);	wait (S);

	signal (S);	signal (Q);
	signal (Q);	signal (S);

- **Starvation** – indefinite blocking. A thread may never get the chance to acquire a lock if the mutex mechanism is not fair.

29

Deadlock Characterization

[Coffman et al. 1971]

Deadlock can arise only if **four conditions** hold simultaneously:

- **Mutual exclusion:** only one thread at a time can use a resource.
- **Hold and wait:** a thread holding at least one resource is waiting to acquire additional resources held by other threads.
- **No preemption of resources:** a resource can be released only voluntarily by the thread holding it, after that thread has completed its task.
- **Circular wait:** there exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting threads such that
 - P_0 is waiting for a resource that is held by P_1 ,
 - P_1 is waiting for a resource that is held by P_2 , ...,
 - P_{n-1} is waiting for a resource that is held by P_n , and
 - P_n is waiting for a resource that is held by P_0 .

30

Coordinating Shared Access (5)



- Must also rely on implementation for efficiency
- Time to lock / unlock mutex or synchronize threads varies widely between different platforms
- A mutex that all threads access serializes the threads!
 - Convoing
 - Goal: Make critical section as short as possible

```
// in gettask():
int tmpindex; // local (thread-private) variable
pthread_mutex_lock( &mutex );
tmpindex = p.index++;
pthread_mutex_unlock( &mutex );
task = p.queue [ tmpindex ];
```

Possibly slow shared memory access now outside critical section

C. Kessler, IDA, Linköping universitet.

31

Coordinating Shared Access (6)



- When programming on this level of abstraction: can minimize serialization, but not avoid
 - Example: Fine-grained locking
- Better: avoid mutex and similar constructs, and use higher-level data structures that are lock-free
 - Example: NOBLE
- Also: Transactional memory

More about this in Lecture 4

C. Kessler, IDA, Linköping universitet.

32

Performance Issues with Threads on Multicores




Performance Issue: Thread Pools

- For a multithreaded process: Create a number of threads in a pool where they await work

- Advantages:
 - Faster to service a request with an existing thread than to create a new thread
 - Allows the number of threads in the application(s) to be bound to the size of the pool

- Win32 API
- OpenMP



TDD868, C. Kessler, IDA, Linköping universitet.

4.34

Performance Issue: Spinlocks on Multiprocessors



- Recall busy waiting at spinlocks:

```
// ... lock initially 0 (unlocked)
while ( ! test_and_set( &lock ) )
;
// ... the critical section ...
lock = 0;
```

- Test_and_set in a tight loop → high bus traffic on multiprocessor

- Cache coherence mechanism must broadcast all writing accesses (incl. t&s) to lock immediately to all writing processors, to maintain a consistent view of lock's value
 - contention
 - degrades performance

Solution 1: TTAS

- Combine with ordinary read:


```
while ( ! test_and_set( &lock ) )
while ( lock )
;
```

// ... the critical section ...

- Most accesses to lock are now reads
 - less contention, as long as lock is not released.

Solution 2: Back-Off

- while (! test_and_set(&lock)) do_nothing_for (short_time);


```
// ... the critical section ...
```

- Exponential / random back-off

35

Performance Issue: Manual Avoidance of Idle Waiting



- Thread that unsuccessfully tried to acquire mutex is blocked but not suspended
 - busy waiting, idle ☹
- Can find out that mutex is locked and do something else:


```
pthread_mutex_trylock ( &mutex_lock );
```

 - If mutex is unlocked, returns 0
 - If mutex is locked, returns EBUSY
- Useful for locks that are not accessed too frequently and for threads having the chance to do something else


36

Performance Issue: Thread Pinning



Programmer may want tight control over thread-to-core mapping of the underlying OS CPU scheduler, e.g. due to:

- Caching, NUMA ("non-uniform memory access [time]")
→ not all cores are equally "close" to all memory locations
→ map thread to a specific core (subset) it has "affinity" to
- Application / runtime system might do its own load balancing (see later)
- Measurement reproducibility problem
 - Measured time, energy can vary considerably if the thread mapping changes unexpectedly



Solution: "pin" a thread to a specific core

- Constrain CPU scheduler to always map it to that same core

37

Thread Pinning with Pthreads



- `#include <pthread.h>`

```
int pthread_setaffinity_np (
    pthread_t thread,
    size_t cpusetsize,
    const cpu_set_t *cpuset
);
```

Bitvector-like data structure describing a subset of the available CPU set):
i th bit = 1 iff core i is in the affinity set of this thread.

```
int pthread_getaffinity_np (
    pthread_t thread,
    size_t cpusetsize,
    cpu_set_t *cpuset
);
```

38

Thread Pinning with Pthreads



- Pthreads provides macros for defining and comparing CPU sets.
- Example: pin the current thread to core i (in 0,...,p-1):

```
#include <sched.h>
#include <pthread.h>
...

cpu_set_t cpuset;

CPU_ZERO(&cpuset); // empty CPU mask
CPU_SET(i, &cpuset); // set i-th "bit" in cpuset

pthread_t this_thread = pthread_self(); // get this thread's handle
pthread_setaffinity_np(this_thread, sizeof(cpu_set_t), &cpuset);
```

Better Programmability for Thread Programming



Short overview of OpenMP™


(see TDDC78 for in-depth treatment of OpenMP)

OpenMP™



- Standard for shared-memory thread programming
- Developed for incremental parallelization of HPC code
- Directives (e.g. `#pragma omp parallel`)
- Support in recent C compilers, e.g. gcc from v.4.3 and later
- High-level constructs for data and work sharing
 - Low-level thread programming still possible

```
#include <omp.h>
...
#pragma omp parallel shared(N) private(i)
{ // creating a team of OMP_NUM_THREADS threads
    ...
    #pragma omp for schedule(static)
    for (i=0; i<N; i++)
        domuchwork(i);
}
```



Work (here: iterations of for loop) shared among all threads of the current team

41

→EXAMPLE

Performance Issue: Load Balancing



- Longest-running process determines parallel execution time ("makespan")
- Minimized by **load balancing**
 - Static – mapping of tasks to cores *before* runtime, no OH
 - Dynamic – mapping done *at* runtime
 - ▶ Shared (critical section) or distributed work pool
 - ▶ On-line problem – don't know the future, only the past
 - Heuristics such as best-fit, random work stealing

Example: Parallel loop, iterations of unknown+varying workload

```
#pragma omp parallel for schedule(dynamic)
for (i=0; i<N; i++) work(i, unknownworkload(i));
```

42




Programming with Tasks (rather than threads)

Example: Cilk

Task-level programming



- Threads are too inflexible, still too coarse-grained, only moderately portable
- Idea: **Program for tasks, not for threads**
- Lots of tasks, with explicit or implicit dependences
 - Created dynamically by the application
 - Execute non-preemptively
 - Maintained and scheduled by a run-time system running (at user level) on top of *worker threads* provided by underlying hardware/OS
 - Scheduling can be driven by operand data-flow
 - Central task pool vs. Work-stealing scheduler
- Examples:
 - Cilk
 - OpenMP 3.0 task model and runtime system
 - StarPU for heterogeneous / hybrid multicores
 - StarSS / OmpSS / SMPSS



44

Cilk



- supertech.csail.mit.edu/cilk
- algorithmic multithreaded language
- programmer specifies parallelism and exploits data locality
- runtime system schedules computation to a parallel platform
- extension of C
- fork-join execution style
 - typ. overhead for spawning on SMP ca. 4x time for subroutine call
- Commercial branch Cilk++ for C++, bought by Intel in 2009
 - Intel Cilk™ Plus, part of Intel Parallel Building Blocks

45

Cilk: Fine-grained multithreading



```
cilk int fib (int n)
{
  if (n < 2) return n;
  else {
    int x, y;
    x = spawn fib (n-1);
    y = spawn fib (n-2);
    sync;
    return (x+y);
  }
}
```

- Expose massive, fine-grained thread-level parallelism (*tasks*)
- Restricted synchronization – no mutual exclusion, possibly non-preemptable execution
- Use low-overhead dynamic (work-stealing) task scheduler for automatic load balancing


→ Raise the abstraction level: program for tasks, not threads

46

Execution of Cilk programs



- Cilk tasks may execute local tasks (black edges = continuation of control flow in execution trace), **spawn** child tasks or **synchronize** with child tasks (blue dependence edges).
- Execution forms a directed acyclic graph (DAG), task graph
- DAG depth = Length of longest path (critical path) is a lower bound on parallel execution time




47

Execution of Cilk Programs (2)




- Execution schedule:**
 - maps tasks to *processors* X *time steps* (Gantt chart)
 - follows the DAG precedence constraints
 - each processor executes at most one task per time step
- A task is ready for scheduling if all its dependencies are saturated.
- Simple dynamic scheduling algorithm: **Greedy (list) scheduling**
 - Maintain central task pool and the set ("front") of ready tasks
 - At each time step issue all (max. *p*) ready tasks
- In practice, decentral work-stealing schedulers work well.



48

Work-Stealing

- Each processor / core maintains its own local task pool
 - A doubly-ended queue (deque)
 - Spawned tasks (frames) pushed on top
 - Processor pops its next task from top
- Like an ordinary local call stack → low overhead unless idle
- Idle processor steals a task from a randomly chosen victim processor's task queue bottom
- If the victim is also idle, try another one...



Work-Stealing – Implementation (1)

```
taskqueue tasks; // one per worker thread

void spawn ( task t )
{
    myworker.tasks.push( t );
}

void sync() // wait for last spawned subtask:
{
    (status, t) = myworker.tasks.pop();
    if (status == STOLEN) {
        while (! t.done)
            ; // wait for t

        return t.result;
    }
    else // have to do it myself:
        return t.execute();
}

void steal_work( victim )
{
    t = victim.tasks.steal();
    if (t != NULL) { // got one:
        t.thief = myworker;
        t.result = t.execute();
        t.done = TRUE;
    }
}

thread worker ( id, root_task )
{
    if (id == 0)
        root_task.execute();
    else
        forever
            steal_work (rand_victim());
}
```



Work-Stealing – Implementation (2)

```
taskqueue tasks; // one per worker thread

void spawn ( task t )
{
    myworker.tasks.push( t );
}

void sync() // wait for last spawned subtask:
{
    (status, t) = myworker.tasks.pop();
    if (status == STOLEN) {
        while (! t.done)
            // try to steal back ("leapfrogging")
            steal_work( t.thief );
        return t.result;
    }
    else // have to do it myself:
        return t.execute();
}

void steal_work( victim )
{
    t = victim.tasks.steal();
    if (t != NULL) { // got one:
        t.thief = myworker;
        t.result = t.execute();
        t.done = TRUE;
    }
}

thread worker ( id, root_task )
{
    if (id == 0)
        root_task.execute();
    else
        forever
            steal_work (rand_victim());
}
```

Why is this a good idea?



Futures

Implicit task creation and synchronization with Futures

Futures

Here: in Java (5+) Concurrency

```
class FibTask implements Callable<Integer>
{
    static ExecutorService exec = Executors.newCachedThreadPool();
    int arg;

    public FibTask( int n )
    {
        arg = n;
    }

    public Integer call()
    {
        if (arg > 2) {
            Future<Integer> left = exec.submit( new FibTask( arg - 1 ) );
            Future<Integer> right = exec.submit( new FibTask( arg - 2 ) );
            return left.get() + right.get();
        }
        else
            return 1;
    }
}
```

A Callable<T> object represents a task that return a value of type T. Method call() (without parameter) returns the result.

Creates a thread pool to which tasks can be submitted, waited for, and killed.

Java thread-pools abstract from platform-specific thread management

spawn two tasks to compute fib(arg-1) and fib(arg-2), with handles left and right, respectively

get(): blocks until future value (arg) has been written or task is completed

See e.g. chapter 16 [Herlihy/Shavit'08]

Futures

```
Future<Integer> left = exec.submit( new FibTask(arg-1));
... = left.get() + ...;
```

- A **future call** by a thread T1 starts a new thread T2 to calculate one or more values and allocates a **future cell** for each of them.
- T1 is passed a read-reference to each future cell and continues immediately.
- T2 is passed a write-reference to each future cell
- Such references can be passed on to other threads
- As (T2) computes results, it writes them to their future cells.
- When any thread touches a future cell via a read-reference, the read stalls until the value has been written.
- A future cell is written only once but can be read many times.
- Used e.g. in Tera-C [Callahan/Smith'90], ML+futures [Blueloch/Reid-Miller'97], StackThreads/MP [Taura et al.'99], Java (5+) Concurrency Package [SUN'04], C++11



Questions?



Further Reading (Selection)

- C. Lin, L. Snyder: *Principles of Parallel Programming*. Addison Wesley, 2008. (general introduction; Pthreads)
- B. Wilkinson, M. Allen: *Parallel Programming, 2e*. Prentice Hall, 2005. (general introduction; pthreads, OpenMP, MPI)
- M. Herlihy, N. Shavit: *The Art of Multiprocessor Programming*. Morgan Kaufmann, 2008. (threads; nonblocking synchronization)
- Chandra, Dagum, Kohr, Maydan, McDonald, Menon: *Parallel Programming in OpenMP*. Morgan Kaufmann, 2001.
- Barbara Chapman et al.: *Using OpenMP - Portable Shared Memory Parallel Programming*. MIT press, 2007.

56



Other references

- OpenMP: www.openmp.org
- NOBLE – Library of non-blocking data structures
www.cse.chalmers.se/research/group/noble/
- Cilk – algorithmic multithreading
supertech.csail.mit.edu/cilk
- StarPU runtime system for heterogeneous multicores
runtime.bordeaux.inria.fr/StarPU
- StarSs / OmpSs www.bsc.es

57