



# Optimization and Parallelization of Sequential Programs

## Lecture 8

Christoph Kessler  
IDA / PELAB  
Linköping University  
Sweden

Christoph Kessler, IDA,  
Linköpings universitet, 2013.

## Outline

Towards (semi-)automatic parallelization of sequential programs

- Data dependence analysis for loops
- Some loop transformations
  - Loop invariant code hoisting, loop unrolling, loop fusion, loop interchange, loop blocking / tiling, scalar expansion
- Static loop parallelization
- Run-time loop parallelization
  - Doacross parallelization
  - Inspector-executor method
- Speculative parallelization (later, if time)
- Auto-tuning (later, if time)

C. Kessler, IDA, Linköpings universitet.

2

TDDD56 Multicore and GPU Programming

## Foundations: Control and Data Dependence



- Consider statements  $S, T$  in a sequential program ( $S=T$  possible)
  - Scope of analysis is typically a function, i.e. intra-procedural analysis
  - Assume that a control flow path  $S \dots T$  is possible
  - Can be done at arbitrary granularity (instructions, operations, statements, compound statements, program regions)
  - Relevant are only the read and write effects on memory (i.e. on program variables) by each operation, and the effect on control flow

- **Control dependence**  $S \rightarrow T$ , if the fact whether  $T$  is executed may depend on  $S$  (e.g. condition)

- Implies that relative execution order  $S \rightarrow T$  must be preserved when restructuring the program
- Mostly obvious from nesting structure in well-structured programs, but more tricky in arbitrary branching code (e.g. assembler code)

Example:

```
S: if (...) {
  ...
T:  ...
  ...
}
```

C. Kessler, IDA, Linköpings universitet.

3

TDDD56 Multicore and GPU Programming

## Foundations: Control and Data Dependence



- **Data dependence**  $S \rightarrow T$ , if statement  $S$  may execute (dynamically) before  $T$  and both may access the same memory location and at least one of these accesses is a write

- Means that execution order "S before T" must be preserved when restructuring the program
- In general, only a conservative over-estimation can be determined statically
- **flow dependence**: (RAW, read-after-write)
  - S may write a location  $z$  that  $T$  may read
- **anti dependence**: (WAR, write-after-read)
  - S may read a location  $x$  that  $T$  may overwrites
- **output dependence**: (WAW, write-after-write)
  - both  $S$  and  $T$  may write the same location

Example:

```
S: z = ... ;
...
T: ... = ..z.. ;
```

(flow dependence)

C. Kessler, IDA, Linköpings universitet.

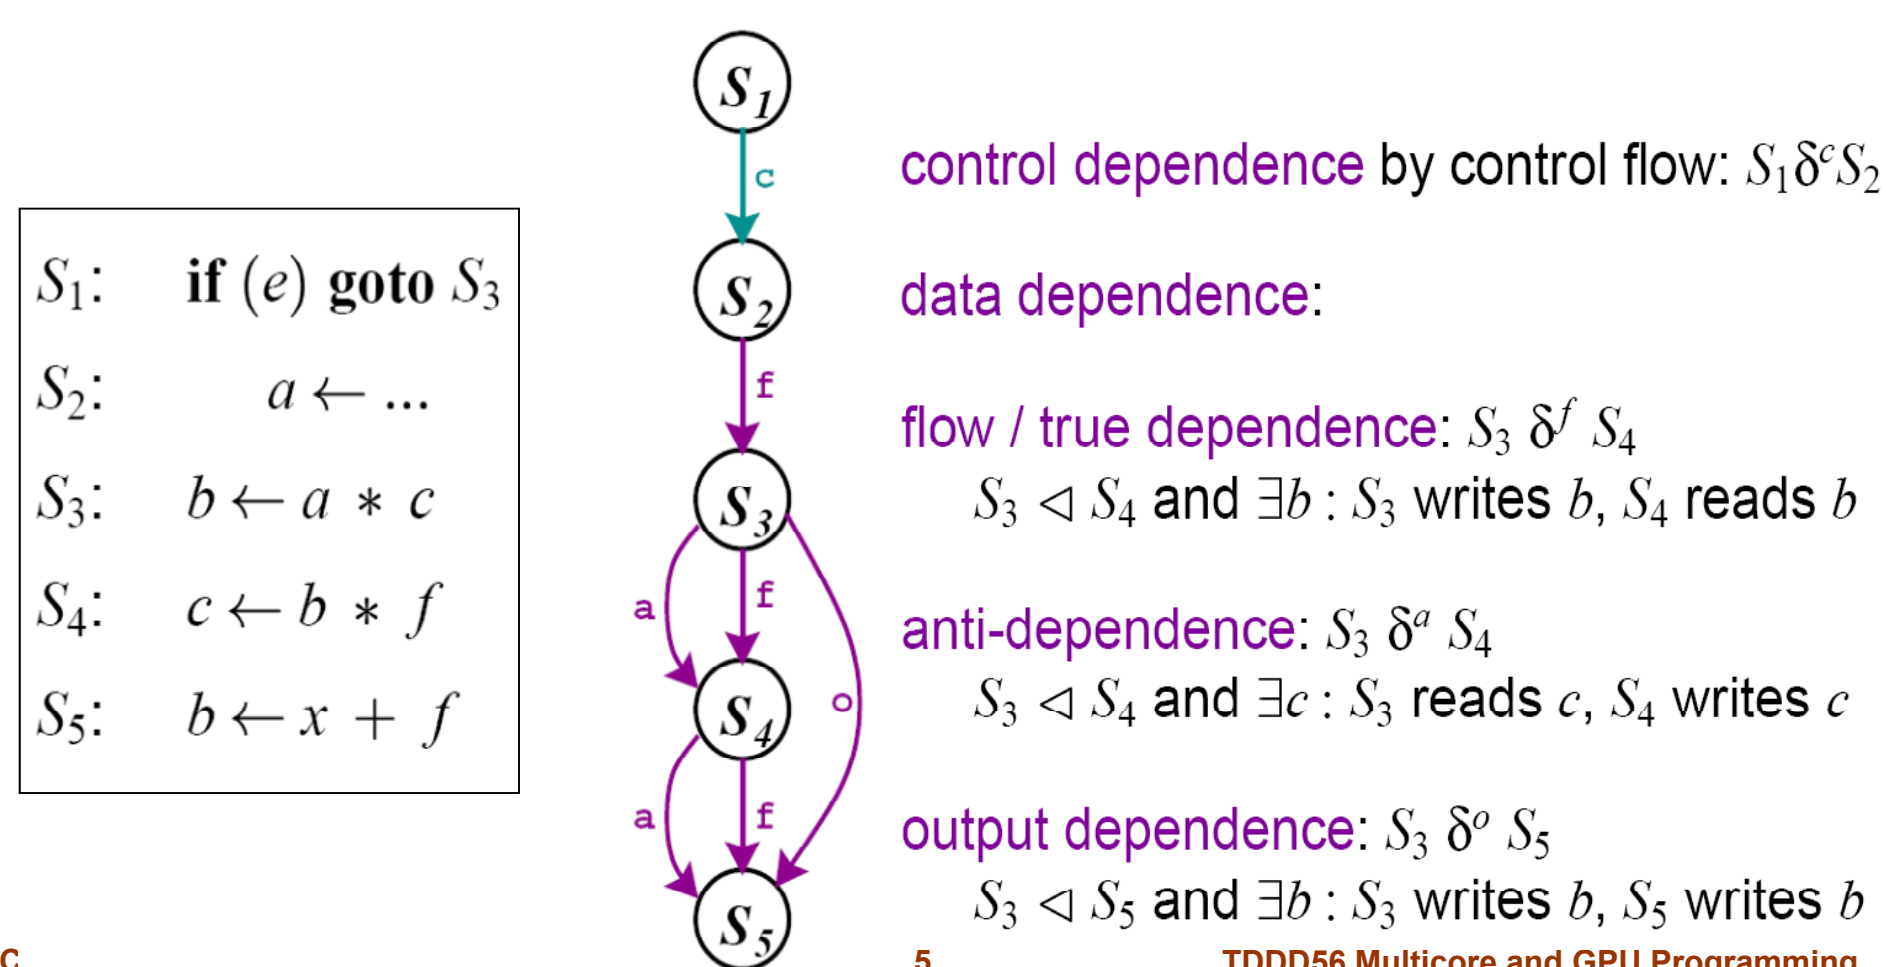
4

TDDD56 Multicore and GPU Programming

## Dependence Graph



- **(Data, Control, Program) Dependence Graph**: Directed graph, consisting of all statements as vertices and all (data, control, any) dependences as edges.



C

5

TDDD56 Multicore and GPU Programming

## Data Dependence Graph



- **Data dependence graph for straight-line code** ("basic block", no branching) is always acyclic, because relative execution order of statements is forward only.

- **Data dependence graph for a loop**:


- Dependence edge  $S \rightarrow T$  if a dependence may exist for some pair of instances (iterations) of  $S, T$
- Cycles possible
- Loop-independent versus loop-carried dependences

Example:

```

for (i=1; i<n; i++) {
S1:  a[i] = b[i] + a[i-1];
S2:  b[i] = a[i];
}
      
```

(assuming we know statically that arrays  $a$  and  $b$  do not intersect)



C. Kessler, II



## Example

```
for i from 2 to 9 do
  S1 X[i] ← Y[i] + Z[i]
  S2 A[i] ← X[i-1] + 1
od
```


(assuming that we statically know that arrays A, X, Y, Z do not intersect, otherwise there might be further dependences)

	$i = 2$	$i = 3$	$i = 4$	...
$S_1$	$X[2] \leftarrow Y[2] + Z[2]$	$X[3] \leftarrow Y[3] + Z[3]$	$X[4] \leftarrow Y[4] + Z[4]$	...
$S_2$	$A[2] \leftarrow X[1] + 1$	$A[3] \leftarrow X[2] + 1$	$A[4] \leftarrow X[3] + 1$	...

There is a loop-carried, forward, flow dependence from  $S_1$  to  $S_2$ .

Iteration space dependence graph:  
(Iterations unrolled)

Data dependence graph:



C. Kessler, IDA, Linköpings universitet.

7

TDDD56 Multicore and GPU Programming

## Why Loop Optimization and Parallelization

Loops are a promising object for program optimizations, including automatic parallelization:

- High execution frequency
  - Most computation done in (inner) loops
  - Even small optimizations can have large impact (cf. Amdahl's Law)
- Regular, repetitive behavior
  - compact description
  - *relatively* simple to analyze statically
- Well researched

C. Kessler, IDA, Linköpings universitet.

8

TDDD56 Multicore and GPU Programming

## Loop Optimizations – General Issues

- Move loop invariant computations out of loops
- Modify the order of iterations or parts thereof

Goals:

- Improve data access locality
- Faster execution
- Reduce loop control overhead
- Enhance possibilities for loop parallelization or vectorization

Only transformations that preserve the program semantics (its input/output behavior) are admissible

- Conservative (static) criterium: preserve data dependences
- Need data dependence analysis for loops (→ DF00100)

C. Kessler, IDA, Linköpings universitet.

9


TDDD56 Multicore and GPU Programming

## Loop Invariant Code Hoisting

- Move loop invariant code out of the loop

- Compilers can do this automatically *if* they can statically find out what code is loop invariant
- Example:

```
for (i=0; i<10; i++)
  a[i] = b[i] + c / d;
```



```
tmp = c / d;
for (i=0; i<10; i++)
  a[i] = b[i] + tmp;
```

C. Kessler, IDA, Linköpings universitet.

10


TDDD56 Multicore and GPU Programming

## Loop Unrolling

- Loop unrolling

- Can be enforced with compiler options e.g. `-funroll=2`
- Example:

```
for (i=0; i<50; i++) {
  a[i] = b[i];
}
```



Unroll  
by 2:

```
for (i=0; i<50; i+=2) {
  a[i] = b[i];
  a[i+1] = b[i+1];
}
```

- ☺ Reduces loop overhead (total # comparisons, branches, increments)
- ☺ Longer loop body may enable further local optimizations (e.g. common subexpression elimination, register allocation, instruction scheduling, using SIMD instructions)
- ⊗ longer code

→ Exercise: Formulate the unrolling rule for statically unknown upper loop limit



## Loop Interchange (1)


- For properly nested loops (statements in innermost loop body only)

- Example 1:


```
for (j=0; j<M; j++)
  for (i=0; i<N; i++)
    a[i][j] = 0.0;
```



```
for (i=0; i<N; i++)
  for (j=0; j<M; j++)
    a[i][j] = 0.0;
```



row-wise  
storage of  
2D-arrays  
in C, Java



- Can improve data access locality in memory hierarchy (fewer cache misses / page faults)

C. Kessler, IDA, Linköpings universitet.

12

TDDD56 Multicore and GPU Programming

## Foundations: Loop-Carried Data Dependences

■ Recall: **Data dependence**  $S \rightarrow T$ , if operation  $S$  may execute (dynamically) before operation  $T$  and both may access the same memory location and at least one of these accesses is a write

$S: z = \dots;$   
 $T: \dots = \dots z \dots;$


- In general, only a conservative over-estimation can be determined statically.

■ Data dependence  $S \rightarrow T$  is called **loop carried** by a loop  $L$  if the data dependence  $S \rightarrow T$  may exist for instances of  $S$  and  $T$  in different iterations of  $L$ .

- Example:

```
L: for (i=1; i<N; i++) {
  Ti: ... = x[i-1];
  Si: x[i] = ...;
}
```

Iteration space:



→ partial order between the operation instances resp. iterations

C. Kessler, IDA, Linköpings universitet.

13


TDDD56 Multicore and GPU Programming

## Loop Interchange (2)

■ Be careful with loop carried data dependences!


- Example 2:

```
for (j=1; j<M; j++)
  for (i=0; i<N; i++)
    a[i][j] = ...a[i+1][j-1]...;
```




```
for (i=0; i<N; i++)
  for (j=1; j<M; j++)
    a[i][j] = ...a[i+1][j-1]...;
```

Iteration space:



Iteration  $(j,i)$  reads location  $a[i+1][j-1]$  that was written in an earlier iteration,  $(i-1,j+1)$



Iteration  $(i,j)$  reads location  $a[i+1][j-1]$ , that will be overwritten in a later iteration  $(i+1,j-1)$

- Interchanging the loop headers would violate the partial iteration order given by the data dependences

C. Kessler, IDA, Linköpings universitet.

14


TDDD56 Multicore and GPU Programming

## Loop Interchange (3)

■ Be careful with loop-carried data dependences!


- Example 3:

```
for (j=1; j<M; j++)
  for (i=1; i<N; i++)
    a[i][j] = ...a[i-1][j-1]...;
```




```
for (i=1; i<N; i++)
  for (j=1; j<M; j++)
    a[i][j] = ...a[i-1][j-1]...;
```

Iteration space:



Iteration  $(j,i)$  reads location  $a[i-1][j-1]$  that was written in earlier iteration  $(j-1,i-1)$



Iteration  $(i,j)$  reads location  $a[i-1][j-1]$  that was written in earlier iteration  $(i-1,j-1)$

- Generally: Interchanging loop headers is only admissible if loop-carried dependences have the same direction for all loops in the loop nest (all directed along or all against the iteration order)

C. Kessler, IDA, Linköpings universitet.

15

TDDD56 Multicore and GPU Programming


## Loop Fusion

■ Merge subsequent loops with same header

- Safe if neither loop carries a (backward) dependence

- Example:

```
for (i=0; i<N; i++)
  a[i] = ...;
  for (i=0; i<N; i++)
    ... = ... a[i] ...;
```



```
for (i=0; i<N; i++) {
  a[i] = ...;
  ... = ... a[i] ...;
}
```

For  $N$  sufficiently large,  $a[j]$  will no longer be in the cache at this time

OK – Read of  $a[j]$  still after write of  $a[j]$ , for all  $i$

- Can improve data access locality and reduces number of branches

C. Kessler, IDA, Linköpings universitet.

16

TDDD56 Multicore and GPU Programming

## Loop Iteration Reordering

A transformation that reorders the iterations of a level- $k$ -loop, without making any other changes, is valid if the loop carries no dependence.

Example:

```
for (i=1; i<n; i++)
  for (j=1; j<m; j++)
    for (k=1; k<r; k++)
      S: a[i][j][k] = ... a[i][j-1][k] ...
```

j-loop carries a dependence, its iteration order must be preserved

(=, <, =)

C. Kessler, IDA, Linköpings universitet.

17

TDDD56 Multicore and GPU Programming

## Loop Parallelization

A transformation that reorders the iterations of a level- $k$ -loop, without making any other changes, is valid if the loop carries no dependence.

Example:

```
for (i=1; i<n; i++)
  for (j=1; j<m; j++)
    for (k=1; k<r; k++)
      S: a[i][j][k] = ... a[i][j-1][k] ...
```


j-loop carries a dependence, its iteration order must be preserved

(=, <, =)

It is valid to convert a sequential loop to a parallel loop if it does not carry a dependence.

Example:

```
for (i=1; i<n; i++)
  S: b[i] = 2 * c[i];
```



```
forall (i, 1, n, p)
  b[i] = 2 * c[i];
```

C. Kessler, IDA, Linköpings universitet.

18

TDDD56 Multicore and GPU Programming



## Remark on Loop Parallelization


- Introducing temporary copies of arrays can remove some antidependences to enable automatic loop parallelization

- Example:

```
for (i=0; i<n; i++)
    a[i] = a[i] + a[i+1];
```

- The loop-carried dependence can be eliminated:

```
for (i=0; i<n; i++)
    aold[i+1] = a[i+1];
for (i=0; i<n; i++)
    a[i] = a[i] + aold[i+1];
```



C. Kessler, IDA, Linköpings universitet.

19

TDDD56 Multicore and GPU Programming

## Strip Mining / Loop Blocking / -Tiling

```
for (i=0; i<n; i++)
    a[i] = b[i] + c[i];
```

↓ loop blocking with block size s

```
for (i1=0; i1<n; i1+=s) // loop over blocks
    for (i2=0; i2<min(n-i1,s); i2++) // loop within blocks
        a[i1+i2] = b[i1+i2] + c[i1+i2];
```

Tiling = blocking in multiple dimensions + loop interchange

Goal: increase locality; support vectorization (vector registers)

Reverse transformation: Loop linearization

C. Kessler, IDA, Linköpings universitet.

20

TDDD56 Multicore and GPU Programming

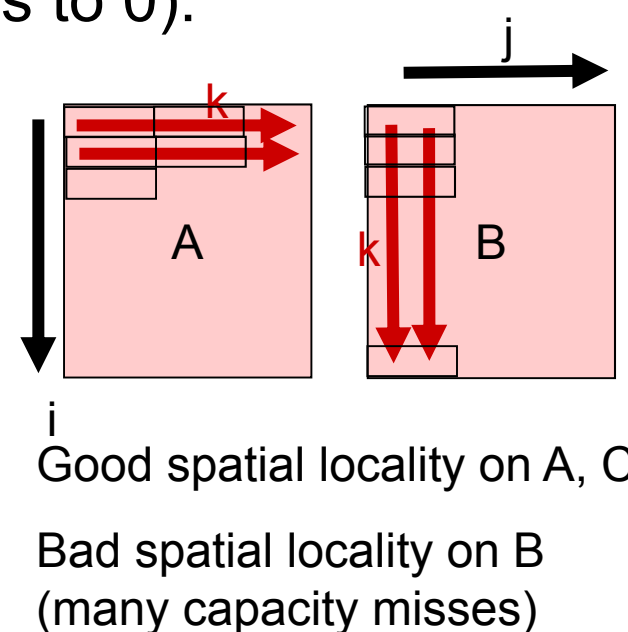
## Tiled Matrix-Matrix Multiplication (1)

- Matrix-Matrix multiplication  $C = A \times B$   
here for square ( $n \times n$ ) matrices  $C, A, B$ , with  $n$  large ( $\sim 10^3$ ):

$$C_{ij} = \sum_{k=1..n} A_{ik} B_{kj} \quad \text{for all } i, j = 1..n$$

- Standard algorithm for Matrix-Matrix multiplication  
(here without the initialization of C-entries to 0):

```
for (i=0; i<n; i++)
    for (j=0; j<n; j++)
        for (k=0; k<n; k++)
            C[i][j] += A[i][k] * B[k][j];
```



C. Kessler, IDA, Linköpings universitet.

21


TDDD56 Multicore and GPU Programming

## Tiled Matrix-Matrix Multiplication (2)

- Block each loop by block size S  
(choose S so that a block of A, B, C fit in cache together).  
then interchange loops

- Code after tiling:

```
for (ii=0; ii<n; ii+=S)
    for (jj=0; jj<n; jj+=S)
        for (i=ii; i < ii+S; i++)
            for (j=jj; j < jj+S; j++)
                for (k=0; k<n; k++)
                    C[i][j] += A[i][k] * B[k][j];
```



Good spatial locality  
for A, B and C

C. Kessler, IDA, Linköpings universitet.

22

TDDD56 Multicore and GPU Programming

## Scalar Expansion / Array Privatization

promote a scalar temporary to an array to break a dependence cycle

```
if N ≥ 1
    allocate t'[1..N]
    for i from 1 to N do
        t'[i] ← a[i] + b[i]
        c[i] ← t'[i] + 1
    od
    t ← t'[N] // if t live on exit
fi
```

expand scalar t:

```
for i from 1 to N do
    t ← a[i] + b[i]
    c[i] ← t + 1
od
```

+ removes the loop-carried antidependence due to t  
→ can now parallelize the loop!

- needs more array space

Loop must be countable, scalar must not have upward exposed uses.

May also be done conceptually only, to enable parallelization:

just create one private copy of t for every processor = **array privatization**

C. Kessler, IDA, Linköpings universitet.

23

TDDD56 Multicore and GPU Programming

## Idiom recognition and algorithm replacement

Traditional loop parallelization fails for loop-carried dep. with distance 1:

```
S0: s = 0;
    for (i=1; i<n; i++)
        s = s + a[i];
S1: a[0] = c[0];
    for (i=1; i<n; i++)
        a[i] = a[i-1] * b[i] + c[i];
```

↓ Idiom recognition (pattern matching)

```
S1': s = VSUM( a[1:n-1], 0 );
```

```
S3': a[0:n-1] = FOLR( b[1:n-1], c[0:n-1], mul, add );
```

↓ Algorithm replacement

```
S1'': s = par_sum( a, 0, n, 0 );
```

C. Kessler, IDA, Linköpings universitet.

24

TDDD56 Multicore and GPU Programming

C. Kessler: Pattern-driven automatic parallelization. *Scientific Programming*, 1996.  
A. Shafiee-Sarvestani, E. Hansson, C. Kessler: Extensible recognition of algorithmic patterns in DSP programs for automatic parallelization. *Int. J. on Parallel Programming*, 2013.





## For further loop transformations...

... see **DF00100** (TDDC86)  
**Advanced Compiler Construction**

Index set splitting, Loop unswitching,  
Loop skewing, Loop distribution,  
Software Pipelining of Loops, ...

Christoph Kessler, IDA,  
Linköpings universitet, 2013.

## Remark on static analyzability (1)



- Static dependence information is always a (safe) overapproximation of the real (run-time) dependences
  - Finding out the real ones exactly is statically undecidable!
  - If in doubt, a dependence must be assumed  
→ may prevent some optimizations or parallelization
- One main reason for imprecision is **aliasing**, i.e. the program may have several ways to refer to the same memory location

- Example: Pointer aliasing

```
void mergesort ( int* a, int n )
{
    ...
    mergesort ( a, n/2 );
    mergesort ( a + n/2, n-n/2 );
    ...
}
```

How could a static analysis tool (e.g., compiler) know that the two recursive calls read and write disjoint subarrays of *a*?

C. Kessler, IDA, Linköpings universitet.

26

TDDD56 Multicore and GPU Programming

## Remark on static analyzability (2)



- Static dependence information is always a (safe) overapproximation of the real (run-time) dependences
  - Finding out the latter exactly is statically undecidable!
  - If in doubt, a dependence must be assumed  
→ may prevent some optimizations or parallelization
- Another reason for imprecision are **statically unknown values** that imply whether a dependence exists or not
  - Example: Unknown dependence distance

```
// value of K statically unknown
for ( i=0; i<N; i++ )
{
    ...
    S: a[i] = a[i] + a[K];
    ...
}
```

Loop-carried dependence if  $K < N$ .  
Otherwise, the loop is parallelizable.

C. Kessler, IDA, Linköpings universitet.

27

TDDD56 Multicore and GPU Programming

## Outlook: Runtime Parallelization



Sometimes parallelizability cannot be decided statically.

```
if is_parallelizable(...)
    forall i in [0..n-1] do // parallel version of the loop
        iteration(i);
    od
else
    for i from 0 to n-1 do // sequential version of the loop
        iteration(i);
    od
fi
```

The runtime dependence test `is_parallelizable(...)` itself may partially run in parallel.

C. Kessler, IDA, Linköpings universitet.

28

TDDD56 Multicore and GPU Programming



## Run-Time Parallelization

Christoph Kessler, IDA,  
Linköpings universitet, 2013.

## Goal of run-time parallelization



- Typical target: **irregular loops**

```
for ( i=0; i<n; i++ )
    a[i] = f ( a[ g(i) ], a[ h(i) ], ... );
```

- Array index expressions *g*, *h*... depend on run-time data
- Iterations cannot be statically proved independent (and not either dependent with distance +1)

- **Principle:**

At runtime, inspect *g*, *h* ... to find out the real dependences and compute a schedule for partially parallel execution

- Can also be combined with speculative parallelization

C. Kessler, IDA, Linköpings universitet.

30

TDDD56 Multicore and GPU Programming



## Overview

- **Run-time parallelization of irregular loops**
  - DOACROSS parallelization
  - Inspector-Executor Technique (shared memory)
  - Inspector-Executor Technique (message passing) \*
  - Privatizing DOALL Test \*
- **Speculative run-time parallelization of irregular loops \***
  - LRPD Test \*
- **General Thread-Level Speculation**
  - Hardware support \*

\* = not covered in this course. See the references.

C. Kessler, IDA, Linköpings universitet.

31


TDDD56 Multicore and GPU Programming

## DOACROSS Parallelization

- Useful if loop-carried dependence distances are unknown, but often > 1
- Allow independent subsequent loop iterations to overlap
- Bilateral synchronization between really-dependent iterations

Example:

```
for ( i=0; i<n; i++)
    a[i] = f ( a[ g(i) ], ... );
```



```
sh float aold[n];
sh flag done[n]; // flag (semaphore) array
forall i in 0..n-1 { // spawn n threads, one per iteration
    done[n] = 0;
    aold[i] = a[i]; // create a copy
}
forall i in 0..n-1 { // spawn n threads, one per iteration
    if (g(i) < i) wait until done[ g(i) ];
    a[i] = f ( a[ g(i) ], ... );
    set( done[i] );
}
else
    a[i] = f ( aold[ g(i) ], ... ); set done[i];
```

C. Kessler, IDA, Linköpings universitet.

32

TDDD56 Multicore and GPU Programming

## Inspector-Executor Technique (1)

- Compiler generates 2 pieces of customized code for such loops:

### Inspector

- calculates values of index expression by simulating whole loop execution
  - typically, based on sequential version of the source loop (some computations could be left out)
- computes implicitly the real iteration dependence graph
- computes a parallel schedule as (greedy) wavefront traversal of the iteration dependence graph in topological order
  - all iterations in same wavefront are independent
  - schedule depth = #wavefronts = critical path length



### Executor

- follows this schedule to execute the loop



C. Kessler, IDA, Linköpings universitet.

33

TDDD56 Multicore and GPU Programming

## Inspector-Executor Technique (2)

### Source loop:

```
for ( i=0; i<n; i++)
    a[i] = f ( a[ g(i) ], a[ h(i) ], ... );
```

### Inspector:

```
int wf[n]; // wavefront indices
int depth = 0;
for (i=0; i<n; i++)
    wf[i] = 0; // init.
for (i=0; i<n; i++) {
    wf[i] = max ( wf[ g(i) ], wf[ h(i) ], ... ) + 1;
    depth = max ( depth, wf[i] );
}
```

- Inspector considers only flow dependences (RAW), anti- and output dependences to be preserved by executor



C. Kessler, IDA, Linköpings universitet.

34

TDDD56 Multicore and GPU Programming

## Inspector-Executor Technique (3)


### Example:

```
for (i=0; i<n; i++)
    a[i] = ... a[ g(i) ] ...;
```

### Executor:

```
float aold[n]; // buffer array
aold[1:n] = a[1:n];
for (w=0; w<depth; w++)
    forall (i, 0, n, #) if (wf[i] == w) {
        a1 = (g(i) < i)? a[g(i)] : aold[g(i)];
        ... // similarly, a2 for h etc.
        a[i] = f ( a1, a2, ... );
    }
```

i	0	1	2	3	4	5
g(i)	2	0	2	1	1	0
wf[i]	0	1	0	2	2	1
g(i)<i?	no	yes	no	yes	yes	yes



C. Kessler, IDA, Linköpings universitet.

35

TDDD56 Multicore and GPU Programming

## Inspector-Executor Technique (4)

**Problem:** Inspector remains sequential – no speedup

### Solution approaches:

- Re-use schedule over subsequent iterations of an outer loop if access pattern does not change
  - amortizes inspector overhead across repeated executions
- Parallelize the inspector using doacross parallelization [Saltz, Mirchandaney'91]
- Parallelize the inspector using sectioning [Leung/Zahorjan'91]
  - compute processor-local wavefronts in parallel, concatenate
  - trade-off schedule quality (depth) vs. inspector speed
  - Parallelize the inspector using bootstrapping [Leung/Z.'91]
  - Start with suboptimal schedule by sectioning, use this to execute the inspector → refined schedule



C. Kessler, IDA, Linköpings universitet.

36

TDDD56 Multicore and GPU Programming



## Questions?

Christoph Kessler, IDA,  
Linköpings universitet, 2013.

## Some references on Dependence Analysis Loop optimizations and Transformations



- H. Zima, B. Chapman: *Supercompilers for Parallel and Vector Computers*. Addison-Wesley / ACM press, 1990.
- M. Wolfe: *High-Performance Compilers for Parallel Computing*. Addison-Wesley, 1996.
- R. Allen, K. Kennedy: *Optimizing Compilers for Modern Architectures*. Morgan Kaufmann, 2002.

Idiom recognition and algorithm replacement:

- C. Kessler: Pattern-driven automatic parallelization. *Scientific Programming* 5:251-274, 1996.
- A. Shafiee-Sarvestani, E. Hansson, C. Kessler: Extensible recognition of algorithmic patterns in DSP programs for automatic parallelization. *Int. J. on Parallel Programming*, 2013.

C. Kessler, IDA, Linköpings universitet.

38

TDDD56 Multicore and GPU Programming

## Some references on run-time parallelization



- R. Cytron: Doacross: Beyond vectorization for multiprocessors. Proc. ICPP-1986
- D. Chen, J. Torrellas, P. Yew: An Efficient Algorithm for the Run-time Parallelization of DO-ACROSS Loops, Proc. IEEE Supercomputing Conf., Nov. 2004, IEEE CS Press, pp. 518-527
- R. Mirchandaney, J. Saltz, R. M. Smith, D. M. Nicol, K. Crowley: Principles of run-time support for parallel processors, Proc. ACM Int. Conf. on Supercomputing, July 1988, pp. 140-152.
- J. Saltz and K. Crowley and R. Mirchandaney and H. Berryman: Runtime Scheduling and Execution of Loops on Message Passing Machines, *Journal on Parallel and Distr. Computing* 8 (1990): 303-312.
- J. Saltz, R. Mirchandaney: The preprocessed doacross loop. Proc. ICPP-1991 Int. Conf. on Parallel Processing.
- S. Leung, J. Zahorian: Improving the performance of run-time parallelization. Proc. ACM PPOPP-1993, pp. 83-91.
- Lawrence Rauchwerger, David Padua: The Privatizing DOALL Test: A Run-Time Technique for DOALL Loop Identification and Array Privatization. Proc. ACM Int. Conf. on Supercomputing, July 1994, pp. 33-45.
- Lawrence Rauchwerger, David Padua: The LRPD Test: Speculative Run-Time Parallelization of Loops with Privatization and Reduction Parallelization. Proc. ACM SIGPLAN PLDI-95, 1995, pp. 218-232.

C. Kessler, IDA, Linköpings universitet.

39

TDDD56 Multicore and GPU Programming